

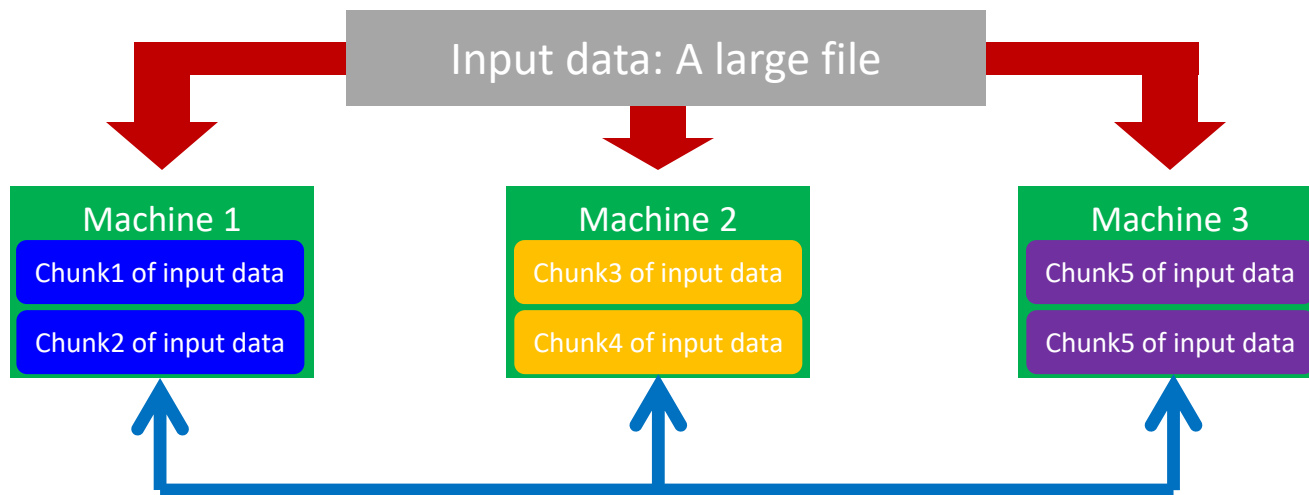
NoSQL Databases

Scaling Traditional Databases

- Traditional RDBMSs can be either scaled:
 - **Vertically** (or **Up**)
 - Can be achieved by hardware upgrades (e.g., faster CPU, more memory, or larger disk)
 - Limited by the amount of CPU, RAM and disk that can be configured on a single machine
 - **Horizontally** (or **Out**)
 - Can be achieved by adding more machines
 - Requires database *sharding* and probably *replication*
 - Limited by the Read-to-Write ratio and communication overhead

Why Sharding Data?

- Data is typically *sharded* (or *striped*) to allow for concurrent/parallel accesses



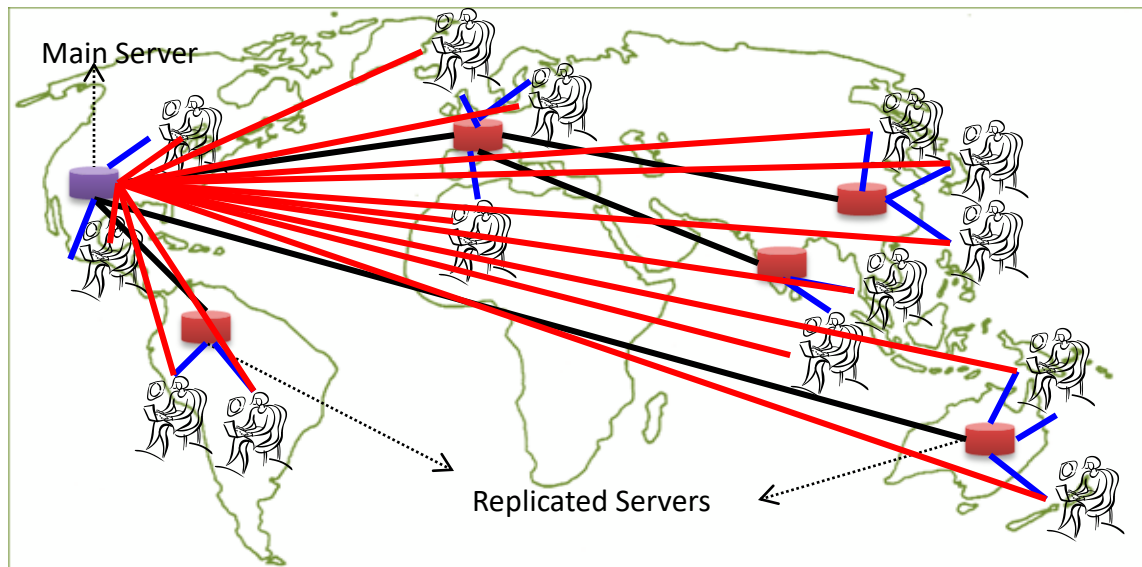
E.g., Chunks 1, 3 and 5 can be accessed in parallel

Why Replicating Data?

- Replicating data across servers helps in:
 - Avoiding performance bottlenecks
 - Avoiding single point of failures
 - And, hence, enhancing scalability and availability

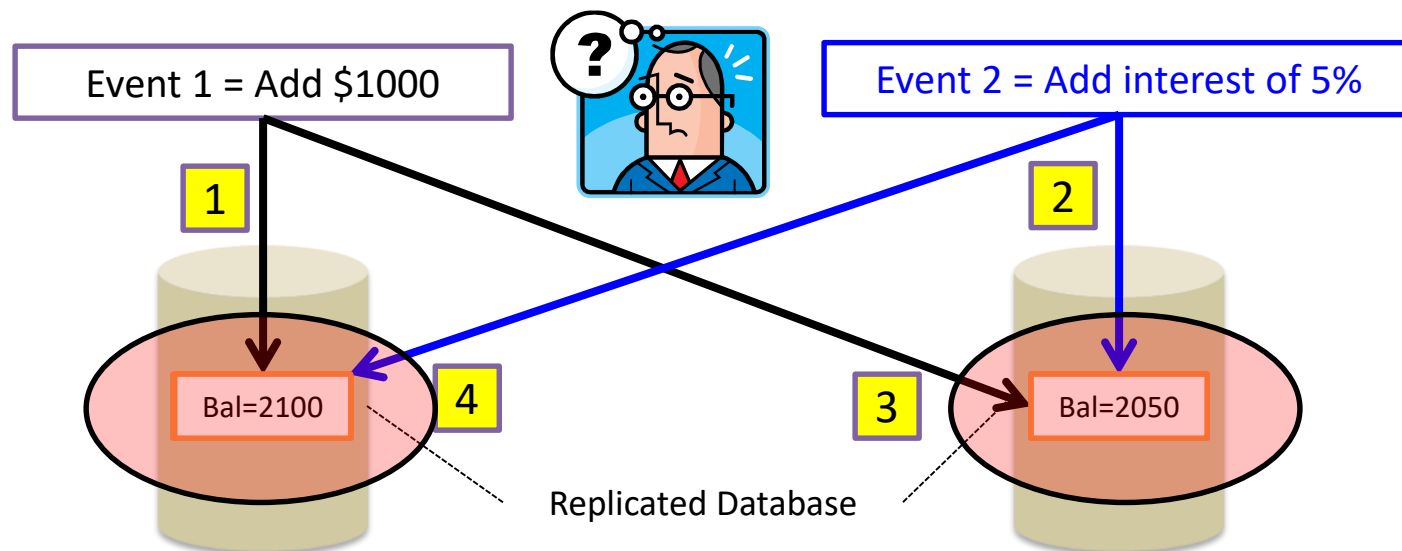
Why Replicating Data?

- Replicating data across servers helps in:
 - Avoiding performance bottlenecks
 - Avoiding single point of failures
 - And, hence, enhancing scalability and availability



But, Consistency Becomes a Challenge

- An example:
 - In an e-commerce application, the bank database has been replicated across two servers
 - Maintaining consistency of replicated data is a challenge




The CAP Theorem

- The limitations of distributed databases can be described in the so called the **CAP theorem**
 - **Consistency**: every node always sees the same data at any given instance (i.e., strict consistency)
 - **Availability**: the system continues to operate, even if nodes in a cluster crash, or some hardware or software parts are down due to upgrades
 - **Partition Tolerance**: the system continues to operate in the presence of network partitions

CAP theorem: any distributed database with shared data, can have at most two of the three desirable properties, C, A or P

The CAP Theorem (*Cont'd*)

- Let us assume two nodes on opposite sides of a network partition:



- Availability + Partition Tolerance forfeit Consistency
- Consistency + Partition Tolerance entails that one side of the partition must act as if it is unavailable, thus forfeiting Availability
- Consistency + Availability is only possible if there is no network partition, thereby forfeiting Partition Tolerance

Large-Scale Databases

- When companies such as Google and Amazon were designing large-scale databases, 24/7 Availability was a key
 - A few minutes of downtime means lost revenue
- When *horizontally* scaling databases to 1000s of machines, the likelihood of a node or a network failure increases tremendously
- Therefore, in order to have strong guarantees on Availability and Partition Tolerance, they had to sacrifice “strict” Consistency (*implied by the CAP theorem*)

The BASE Properties

- The CAP theorem proves that it is impossible to guarantee strict Consistency and Availability while being able to tolerate network partitions
- This resulted in databases with relaxed ACID guarantees
- In particular, such databases apply the BASE properties:
 - Basically Available: the system guarantees Availability
 - Soft-State: the state of the system may change over time
 - Eventual Consistency: the system will *eventually* become consistent

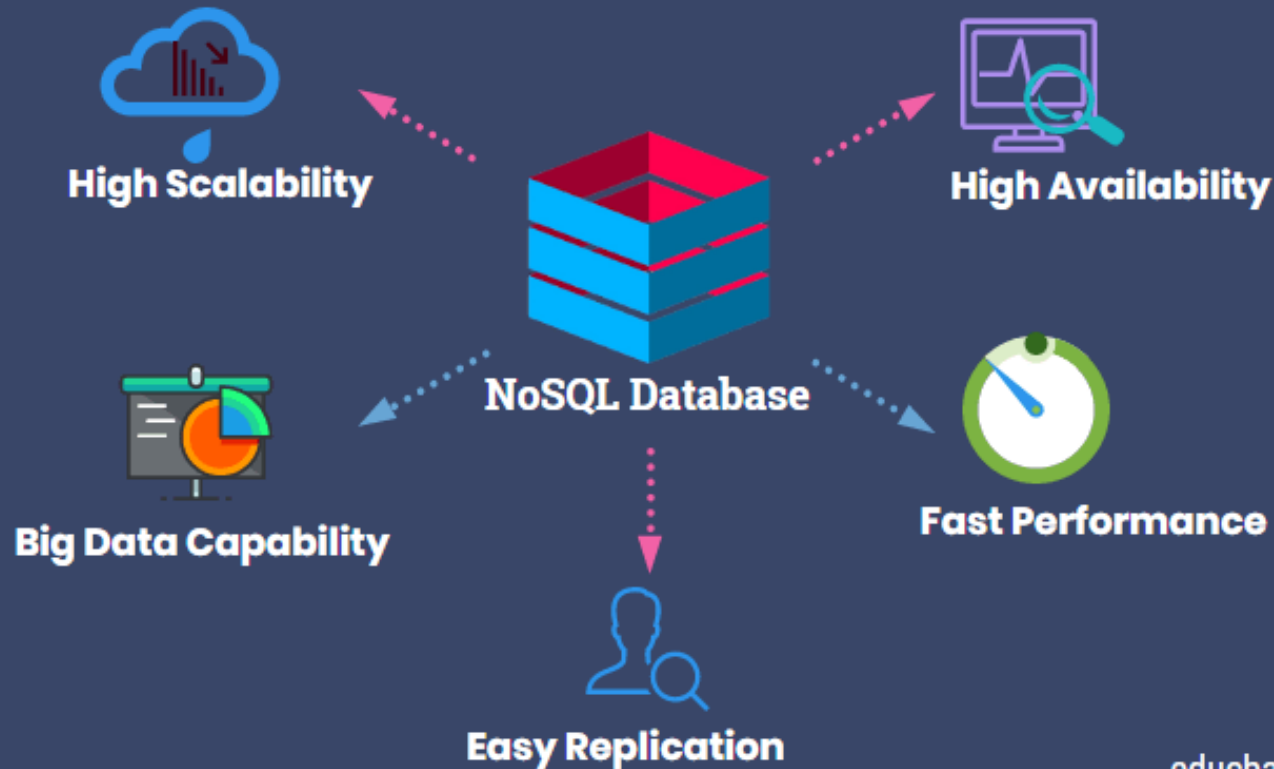
Eventual Consistency

- A database is termed as *Eventually Consistent* if:
 - All replicas will *gradually* become consistent in the absence of updates

NoSQL Databases

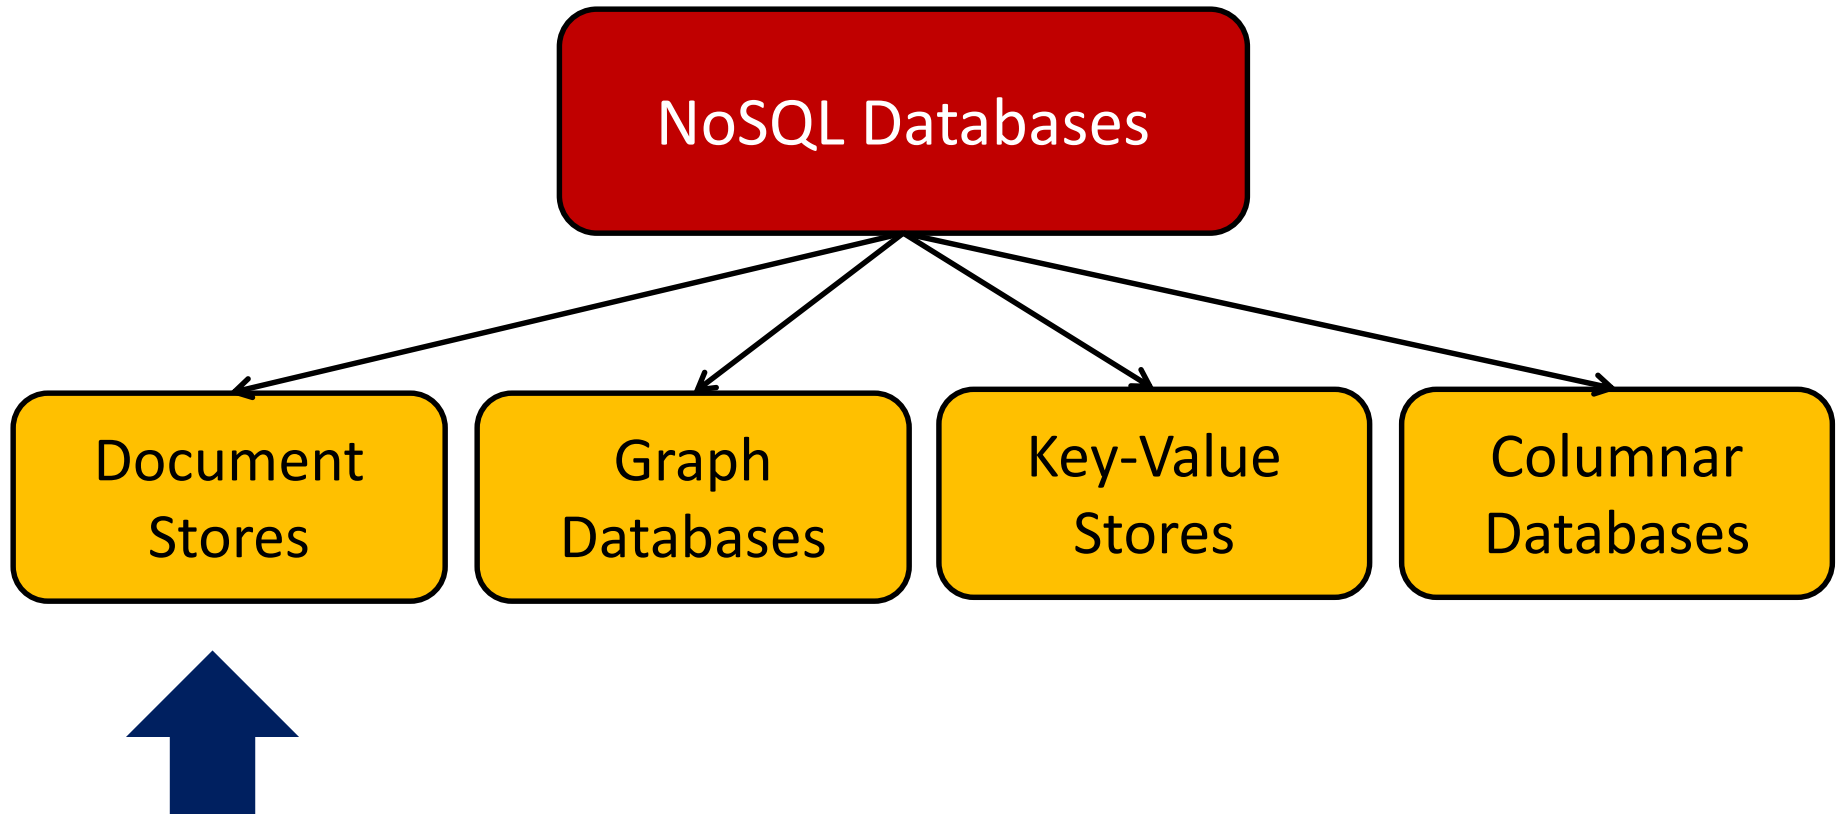
- To this end, a new class of databases emerged, which mainly follow the BASE properties
 - These were dubbed as NoSQL databases
 - E.g., Amazon's Dynamo and Google's Bigtable
- Main characteristics of NoSQL databases include:
 - No strict schema requirements
 - No strict adherence to ACID properties
 - Consistency is traded in favor of Availability

What is NoSQL Database



Types of NoSQL Databases

- Here is a limited taxonomy of NoSQL databases:



Document Stores

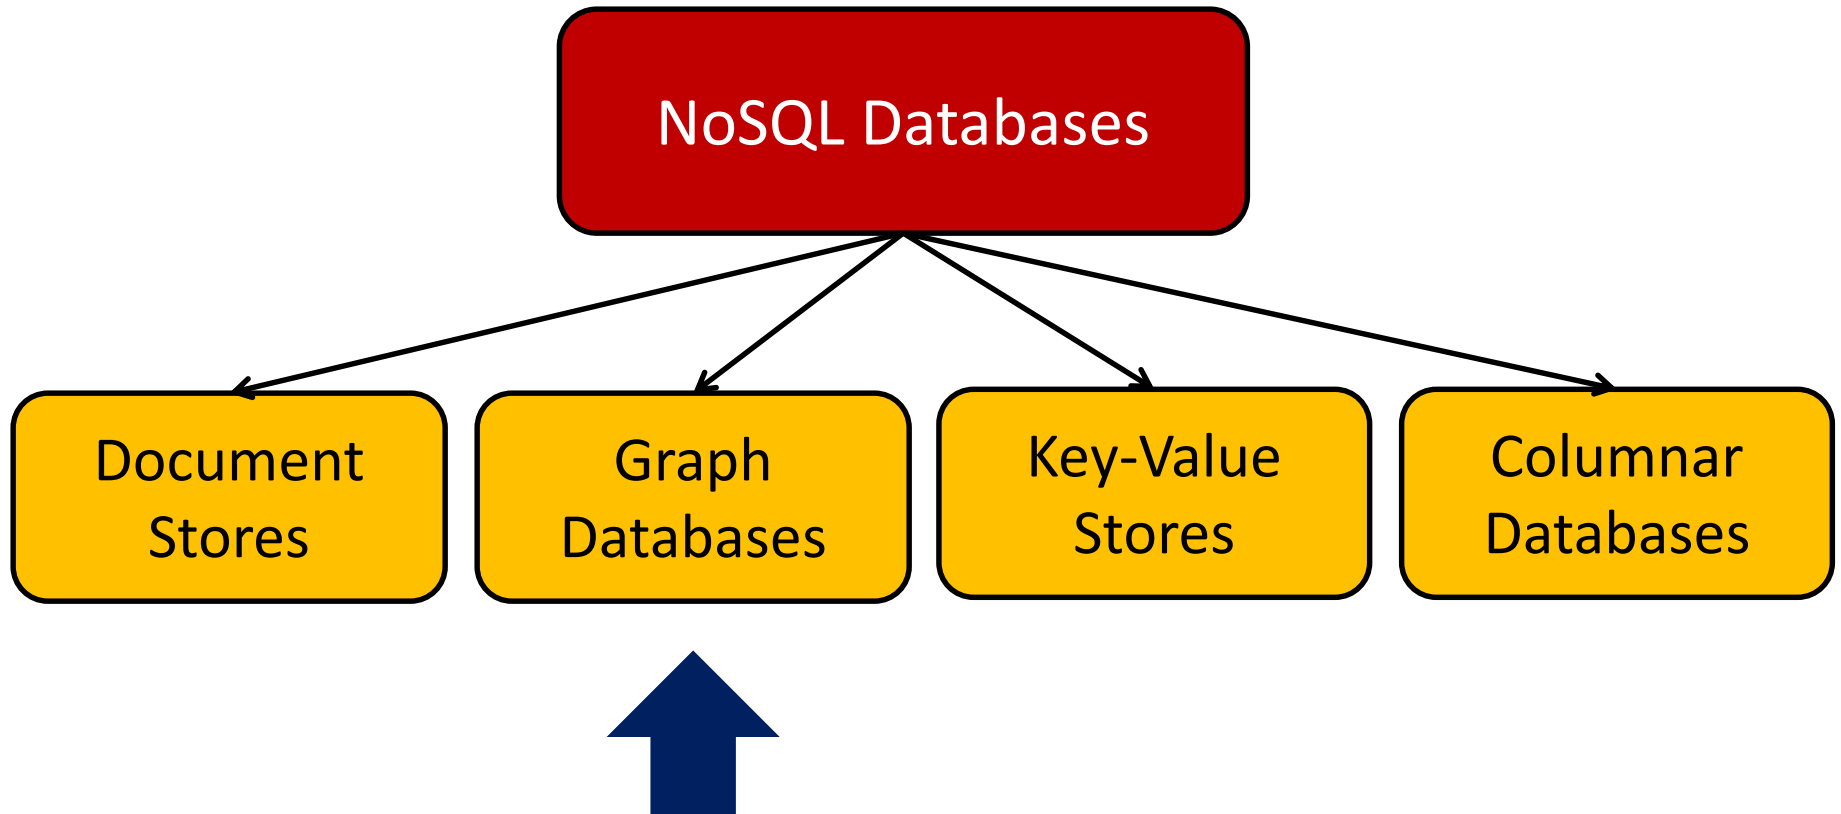
- Documents are stored in some standard format or encoding (e.g., XML, JSON, PDF or Office Documents)
 - These are typically referred to as Binary Large Objects (BLOBs)
- Documents can be indexed
 - This allows document stores to outperform traditional file systems
- E.g., MongoDB and CouchDB (both can be queried using MapReduce)

Document Databases, JSON

```
{  
  _id: ObjectId("51156a1e056d6f966f268f81"),  
  type: "Article",  
  author: "Derick Rethans",  
  title: "Introduction to Document Databases with MongoDB",  
  date: ISODate("2013-04-24T16:26:31.911Z"),  
  body: "This arti..."  
},  
{  
  _id: ObjectId("51156a1e056d6f966f268f82"),  
  type: "Book",  
  author: "Derick Rethans",  
  title: "php|architect's Guide to Date and Time Programming with PHP",  
  isbn: "978-0-9738621-5-7"  
}
```

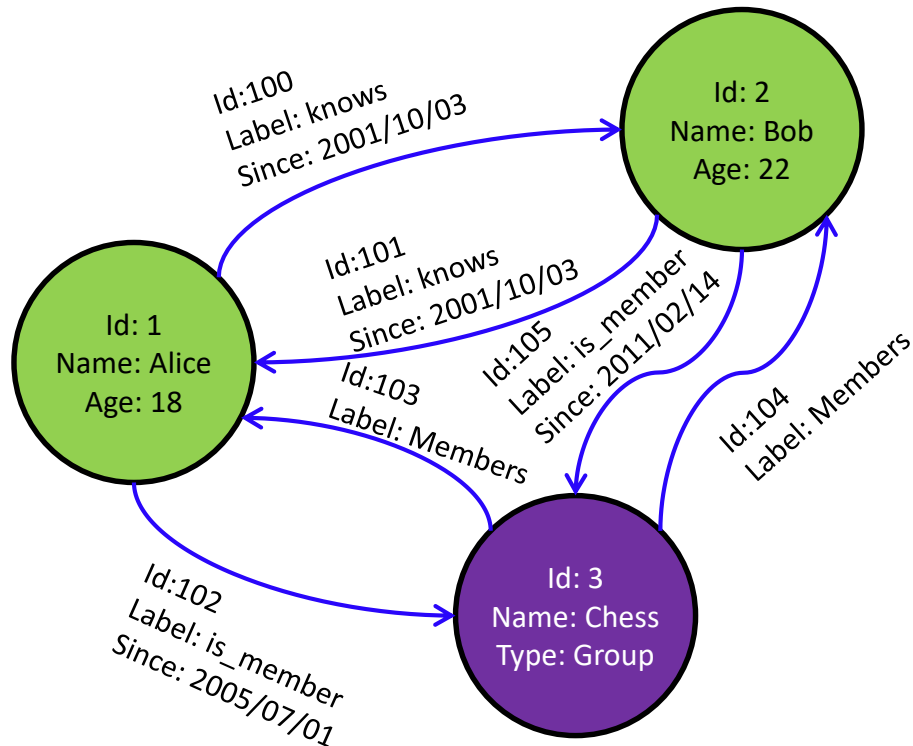

Types of NoSQL Databases

- Here is a limited taxonomy of NoSQL databases:



Graph Databases

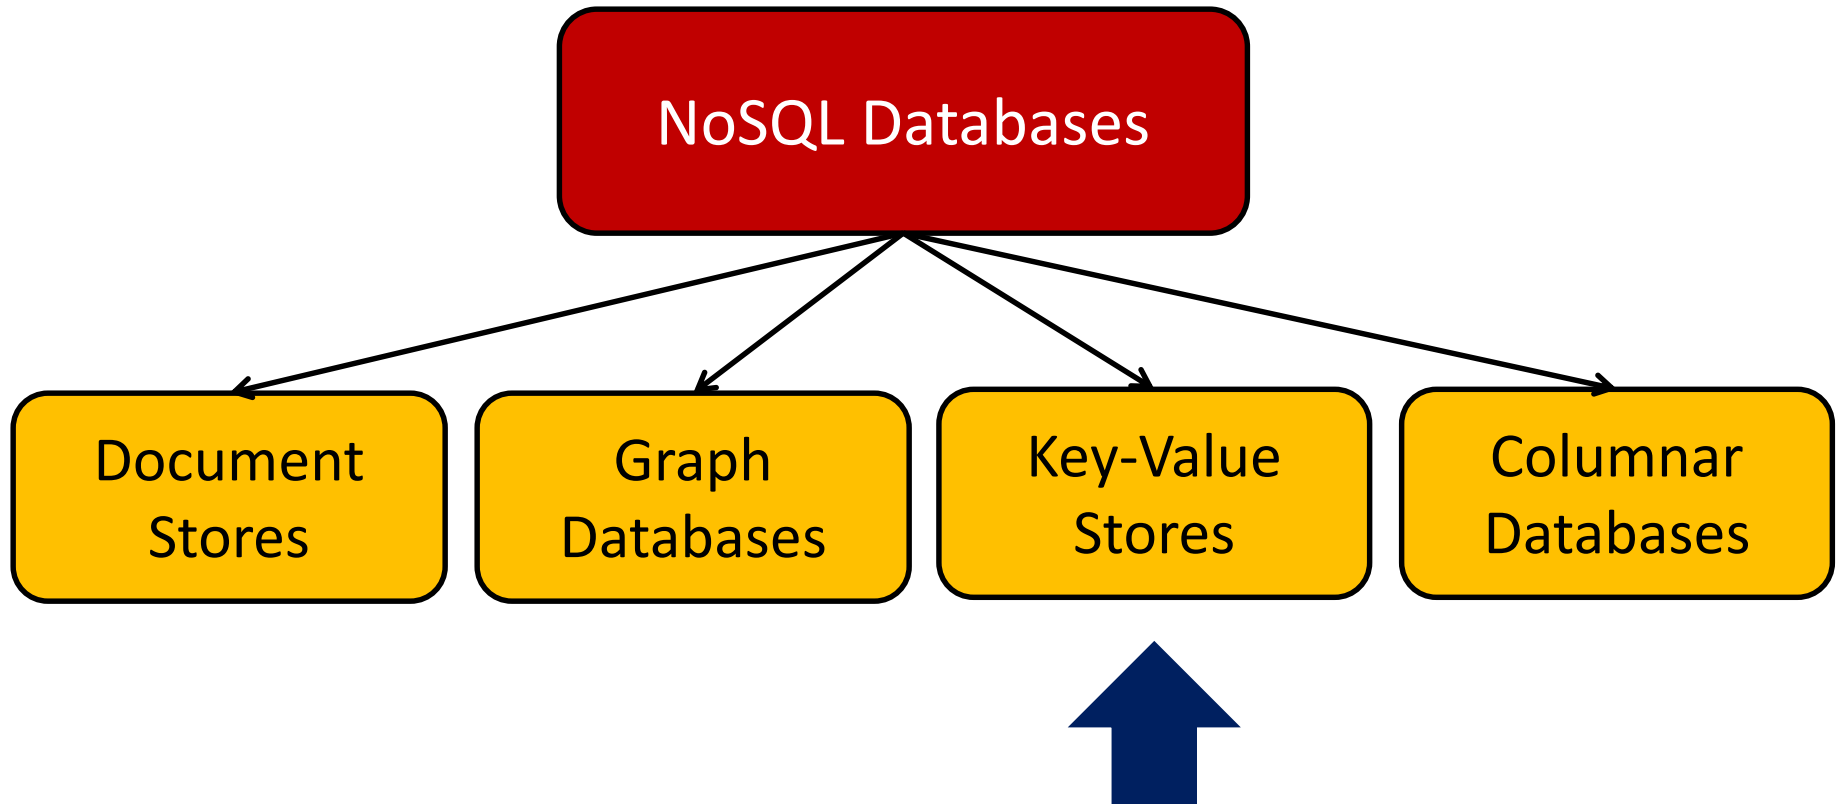
- Data are represented as vertices and edges



- Graph databases are powerful for graph-like queries (e.g., find the shortest path between two elements)
- E.g., Neo4j and VertexDB

Types of NoSQL Databases

- Here is a limited taxonomy of NoSQL databases:

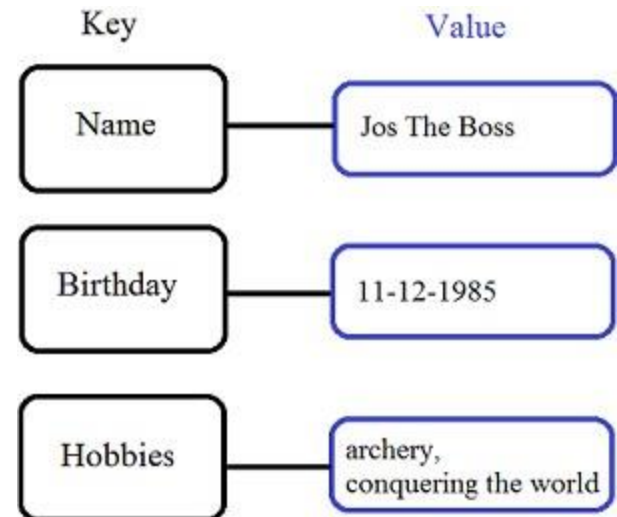


Key-Value Stores

- Keys are mapped to (possibly) more complex value (e.g., lists)
- Keys can be stored in a hash table and can be distributed easily
- Such stores typically support regular CRUD (create, read, update, and delete) operations
 - That is, no joins and aggregate functions

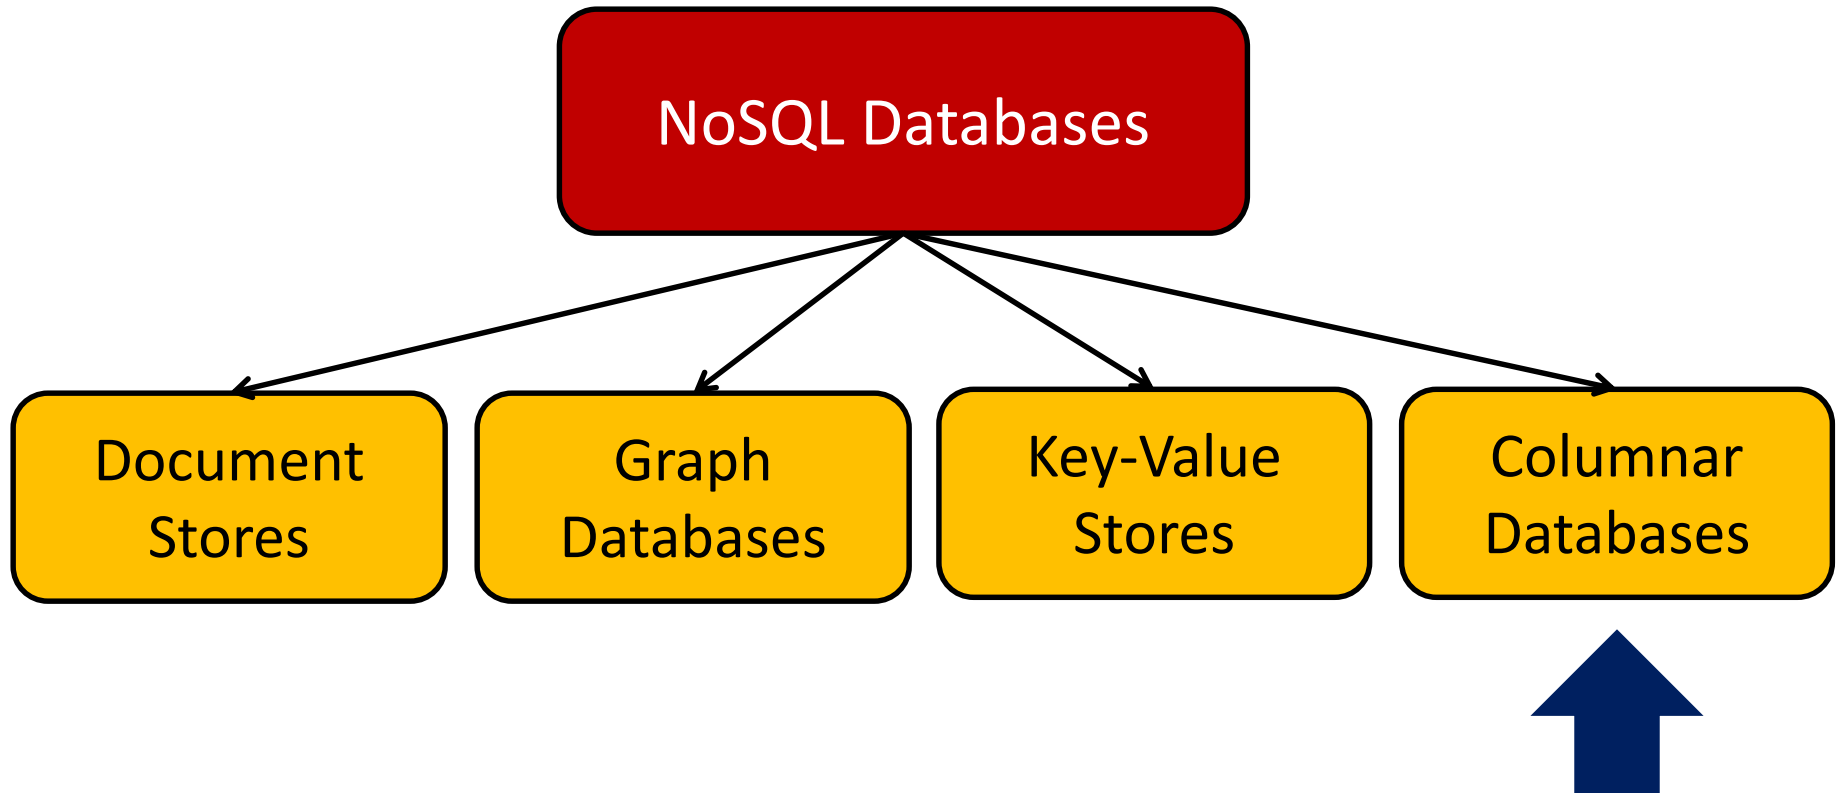
Key/Value stores

- Store data in a schema-less way
- Store data as maps
 - Provide a very efficient average running time algorithm for accessing data
- Notable for:
 - Couchbase
 - Redis
 - Amazon Dynamodb



Types of NoSQL Databases

- Here is a limited taxonomy of NoSQL databases:



Column-Oriented Stores

- Data are stored in a column-oriented way
 - Data efficiently stored
 - Avoids consuming space for storing nulls
 - Columns are grouped in column-families
 - Data isn't stored as a single table but is stored by column families
 - Unit of data is a set of key/value pairs
 - Identified by "row-key"
 - Ordered and sorted based on row-key
- Notable for:
 - Google's Bigtable (used in all Google's services)
 - HBase (Facebook, StumbleUpon, Hulu, Yahoo!, ...)

The diagram illustrates a column-oriented data store structure. It shows a table with the following data:

| Row Key | Students | | Branch | |
|-----------|----------|-----|--------|-----|
| StudentID | Name | Age | Bname | GPA |
| 100 | Ram | 18 | CSE | 7.9 |
| 101 | Sham | 17 | ECE | 8 |
| 102 | John | 18 | EEE | 7.5 |
| 103 | Sam | 17 | CSE | 8.5 |

Annotations in the diagram:

- Column:** A green line points to the rightmost column (GPA).
- Cells:** A red line points to the cell containing the value '8' for StudentID 101.
- Row Key:** A green line points to the 'StudentID' column header.
- Column Families:** Yellow lines point to the 'Students' and 'Branch' column groups.

Where would I use it?

- Where would I use a NoSQL database?
- Do you have somewhere a large set of uncontrolled, unstructured, data that you are trying to fit into a RDBMS?
 - Log Analysis
 - Social Networking Feeds (many firms hooked in through Facebook or Twitter)
 - External feeds from partners
 - Data that is not easily analyzed in a RDBMS such as time-based data
 - Large data feeds that need to be massaged before entry into an RDBMS

Summary

- *NoSQL* (or *Not-Only-SQL*) databases follow the *BASE* properties:
 - Basically Available
 - Soft-State
 - Eventual Consistency
- NoSQL databases have different types:
 - Document Stores
 - Graph Databases
 - Key-Value Stores
 - Columnar Databases